

```
Logging initialized using configuration in file:/
etc/hive/2.3.4.0-3485/0/hive-log4j.properties
hive (default)> SET hive.execution.engine=mr;
hive (default)> create database iaminusa19828041
> ;
OK
Time taken: 1.072 seconds
```

```
hive (default)> use iaminusa19828041;
OK
Time taken: 0.029 seconds
```

```
hive (iaminusa19828041)> create table
tel_cdr_raw(session_id string,simcard_id
string,phone_number string,call_time
string,duration int,cell_id string, network_type
string, drop_reason string) row fo
rmat delimited fields terminated by ',' stored as
textfile;
OK
Time taken: 0.085 seconds
hive (iaminusa19828041)> LOAD data inpath '/user/
iaminusa19828041/simplecdr' into table
tel_cdr_raw;
Loading data to table iaminusa19828041.tel_cdr_raw
Table iaminusa19828041.tel_cdr_raw stats:
[numFiles=1, totalSize=18171]
OK
Time taken: 0.326 seconds
```

```
hive (iaminusa19828041)> select * from tel_cdr_raw
limit 5;
OK
rnd-dcad1bbb-cd4c-4927-929d-94f6a929f061
SIM-00002      PHONE-00009      27/04/2015 13:35:11
17561   cell F   2      REASON 2
rnd-588ff3de-84b6-4903-a1de-e1cd2faca9a8
SIM-00010      PHONE-00004      27/04/2015 13:35:12
17361   cell A   1      REASON 8
rnd-ef5af2cd-51fa-46dd-80b4-c7c089ca67d3
SIM-00008      PHONE-00004      27/04/2015 13:35:13
7061    cell G   2      REASON 9
rnd-169de378-87b4-4c54-9315-7c1991bb1651
SIM-00006      PHONE-00010      27/04/2015 13:35:14
18361   cell A   1      REASON 4
rnd-683c335a-b2b8-4b35-8758-2bff7f1f31ad
SIM-00002      PHONE-00001      27/04/2015 13:35:16
15361   cell E   1      REASON 10
Time taken: 0.095 seconds, Fetched: 5 row(s)
hive (iaminusa19828041)>
```

```
hive (iaminusa19828041)> create table
tel_cdr_proc(session_id string,simcard_id
string,phone_number string,call_time
string,duration int,cell_id string, network_type
string, drop_reason string,call_d
rop_status string);
OK
```

```
hive (iaminusa19828041)> insert into tel_cdr_proc
```

```

> select
>
t2.session_id,t2.simcard_id,t2.phone_number,t2.call
_time,t2.duration,t2.cell_id,t2.network_type,t2.dro
p_reason,
> CASE
> WHEN previous_timestamp is
null then 'NA'
> WHEN ((previous_timestamp)
+ previous_duration - (timestampx)) > 60 then 'NOT-
DROPPED'
> WHEN ((previous_timestamp)
+ previous_duration - (timestampx)) < 60 then
'DROPPED'
> END as call_drop_status
> FROM (
> SELECT
> t1.*,
> LAG(call_time, 1) OVER
(PARTITION BY simcard_id, phone_number ORDER BY
timestampx ) as previous_callstarttime,
> LAG(duration, 1) OVER
(PARTITION BY simcard_id, phone_number ORDER BY
timestampx ) as previous_duration,
> LAG(drop_reason, 1) OVER
(PARTITION BY simcard_id, phone_number ORDER BY
timestampx ) as previous_dropreason,
> LAG(timestampx, 1) OVER
(PARTITION BY simcard_id, phone_number ORDER BY
timestampx ) as previous_timestamp
> FROM (
> SELECT cdr.*,
unix_timestamp(call_time,'dd/MM/yyyy HH:mm:ss') as
timestampx
> FROM tel_cdr_raw cdr
> )t1

```

```
> )t2;
```

```
Query ID = [REDACTED]  
iaminusa19828041_20170729162949_cb689898-220c-4bd8-  
8f07-98a1dfe77c54 [REDACTED]  
Total jobs = 1 [REDACTED]  
Launching Job 1 out of 1 [REDACTED]  
Number of reduce tasks not specified. Estimated  
from input data size: 1 [REDACTED]  
In order to change the average load for a reducer  
(in bytes): [REDACTED]  
    set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
    set hive.exec.reducers.max=<number> [REDACTED]  
In order to set a constant number of reducers:  
    set mapreduce.job.reduces=<number> [REDACTED]  
Starting Job = job_1498404677707_6990, Tracking URL  
= http://a.cloudxlab.com:8088/proxy/  
application_1498404677707_6990/ [REDACTED]  
Kill Command = /usr/hdp/2.3.4.0-3485/hadoop/bin/  
hadoop job -kill job_1498404677707_6990 [REDACTED]  
Hadoop job information for Stage-1: number of  
mappers: 1; number of reducers: 1 [REDACTED]  
2017-07-29 16:29:56,370 Stage-1 map = 0%, reduce =  
0% [REDACTED]  
2017-07-29 16:30:02,684 Stage-1 map = 100%, reduce  
= 0%, Cumulative CPU 3.84 sec [REDACTED]
```

```
Vhive (iaminusa19828041)> select * from  
tel_cdr_proc limit 20;  
OK [REDACTED]  
rnd-3057f491-bf60-49a7-955e-e348a7968236 [REDACTED]  
SIM-00001          PHONE-00006          27/04/2015 13:35:32 [REDACTED]  
461          cell F 2          REASON 7          NA [REDACTED]
```

rnd-28870fcb-71df-416e-adad-c33109baf28e  
SIM-00001 PHONE-00006 27/04/2015 13:35:50  
7061 cell I 2 REASON 5 NOT-DROPPED  
rnd-8def3194-84c2-4489-bc0b-4af10bef73ed  
SIM-00001 PHONE-00007 27/04/2015 13:36:11  
22661 cell B 2 REASON 2 NA  
change-408339cb-8ca9-412d-95a3-a87d87a8c03f  
SIM-00001 PHONE-00007 27/04/2015 13:37:20  
238 cell B 1 REASON 10 NOT-DROPPED  
change-408339cb-8ca9-412d-95a3-a87d87a8c03f  
SIM-00001 PHONE-00007 27/04/2015 13:37:20  
238 cell B 2 REASON 10 NOT-DROPPED  
rnd-8c81b5b9-301f-41bf-907a-48a643e84e79  
SIM-00001 PHONE-00008 27/04/2015 13:35:26  
8561 cell D 1 REASON 7 NA  
rnd-b43aa056-a4c6-47d1-aa6f-2f6ad68ba6ff  
SIM-00001 PHONE-00008 27/04/2015 13:36:14  
7861 cell F 2 REASON 2 NOT-DROPPED  
rnd-f2354458-2551-4b49-8b54-2cc64a00749e  
SIM-00001 PHONE-00009 27/04/2015 13:36:20  
27361 cell G 1 REASON 9 NA  
rnd-d1e6ea89-a7e6-44c6-88f0-e7d6f98edc55  
SIM-00001 PHONE-00010 27/04/2015 13:35:24  
16161 cell F 1 REASON 10 NA  
drop-a00f8eea-5416-4aee-b6a1-e63c8724b352  
SIM-00001-20 PHONE-00009-20 27/04/2015 13:36:47  
2 cell H 1 REASON 4 NA  
drop-a00f8eea-5416-4aee-b6a1-e63c8724b352  
SIM-00001-20 PHONE-00009-20 27/04/2015 13:36:51  
178 cell H 1 REASON 4 DROPPED  
drop-76a06907-ed97-4dd2-9d90-244b14e4e762  
SIM-00001-26 PHONE-00001-26 27/04/2015 13:36:54  
0 cell B 2 REASON 6 NA  
drop-76a06907-ed97-4dd2-9d90-244b14e4e762  
SIM-00001-26 PHONE-00001-26 27/04/2015 13:36:54  
64 cell B 2 REASON 6 DROPPED

```

drop-852dea68-9061-495e-a933-567cde4f8d3e
SIM-00001-9      PHONE-00003-9      27/04/2015 13:36:34
1      cell F 1      REASON 7      NA
drop-852dea68-9061-495e-a933-567cde4f8d3e
SIM-00001-9      PHONE-00003-9      27/04/2015 13:36:37
33      cell F 1      REASON 7      DROPPED
rnd-683c335a-b2b8-4b35-8758-2bffa7f1f31ad
SIM-00002      PHONE-00001      27/04/2015 13:35:16
15361      cell E 1      REASON 10      NA
rnd-8c722bec-8765-41c5-a41f-c6c9cfb30c20
SIM-00002      PHONE-00002      27/04/2015 13:35:42
21061      cell C 2      REASON 4      NA
change-ac5cf3db-d243-4f99-8690-f7d48b8de888
SIM-00002      PHONE-00003      27/04/2015 13:37:24
288      cell E 1      REASON 9      NA
change-ac5cf3db-d243-4f99-8690-f7d48b8de888
SIM-00002      PHONE-00003      27/04/2015 13:37:24
288      cell C 2      REASON 9      NOT-DROPPED
rnd-755150e7-69e5-44f3-b925-67469ef087a6
SIM-00002      PHONE-00005      27/04/2015 13:36:17
11761      cell B 1      REASON 3      NA
Time taken: 0.08 seconds, Fetched: 20 row(s)

```

```

hive (iaminusa19828041)> select
cell_id,drop_reason,count(*) as COT from
tel_cdr_proc where call_drop_status = 'DROPPED'
group by cell_id,drop_
reason order by COT desc limit 10;
Query ID =
iaminusa19828041_20170729163649_e9f12fce-57dc-4e22-
bf3b-52693c7b3ff0
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated
from input data size: 1

```

In order to change the average load for a reducer (in bytes):

set hive.exec.reducers.bytes.per.reducer=<number>

In order to limit the maximum number of reducers:

set hive.exec.reducers.max=<number>

In order to set a constant number of reducers:

set mapreduce.job.reduces=<number>

Starting Job = job\_1498404677707\_6992, Tracking URL

= [http://a.cloudxlab.com:8088/proxy/application\\_1498404677707\\_6992/](http://a.cloudxlab.com:8088/proxy/application_1498404677707_6992/)

Kill Command = /usr/hdp/2.3.4.0-3485/hadoop/bin/hadoop job -kill job\_1498404677707\_6992

Hadoop job information for Stage-1: number of

```
2017-07-29 16:36:55,715 Stage-1 map = 0%, reduce = 0%
2017-07-29 16:37:01,894 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.91 sec
```







